



WHAT'S NEW AT SNOWFLAKE

July, 2026

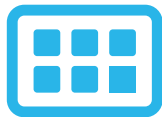
Agentic AI is rewriting how enterprises operate



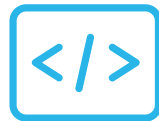
What autonomous AI requires at enterprise scale



**Governed Data
Foundation**



**Model
Choice**



**Universal
Interoperability**



**Agentic
Control Plane**



AI Data Cloud: the agentic foundation

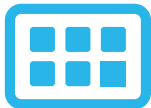


AI DATA CLOUD



Governed Data Foundation

One governed
source of truth
for AI



Model Choice

Choose any model,
any provider,
no lock-in



Universal Interoperability

Connect to any data
or system where
work happens



Agentic Control Plane

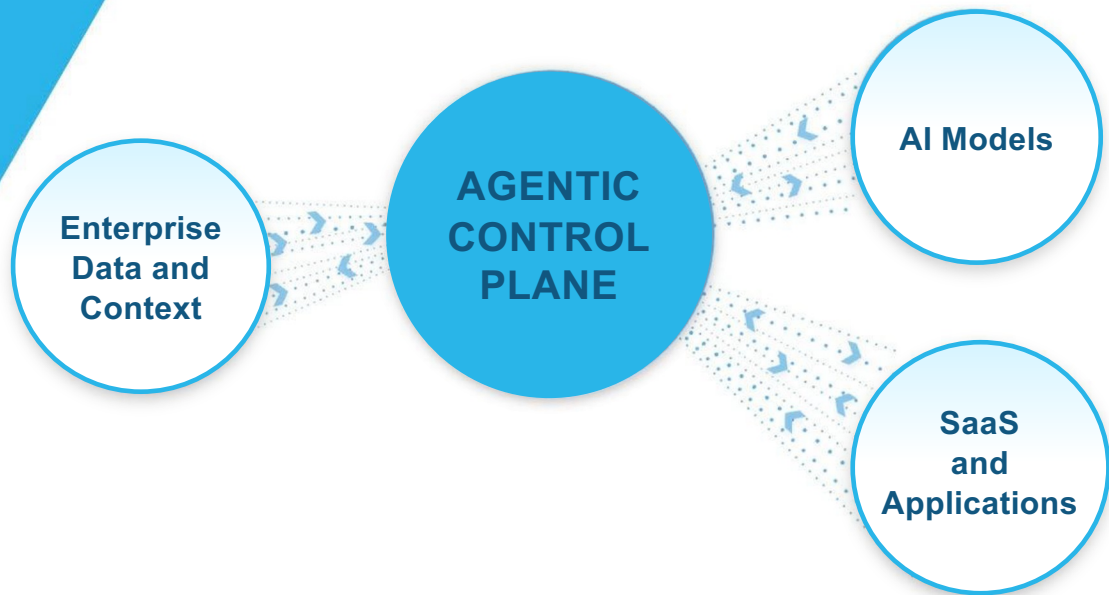
Build, deploy, and
govern agents in
one place



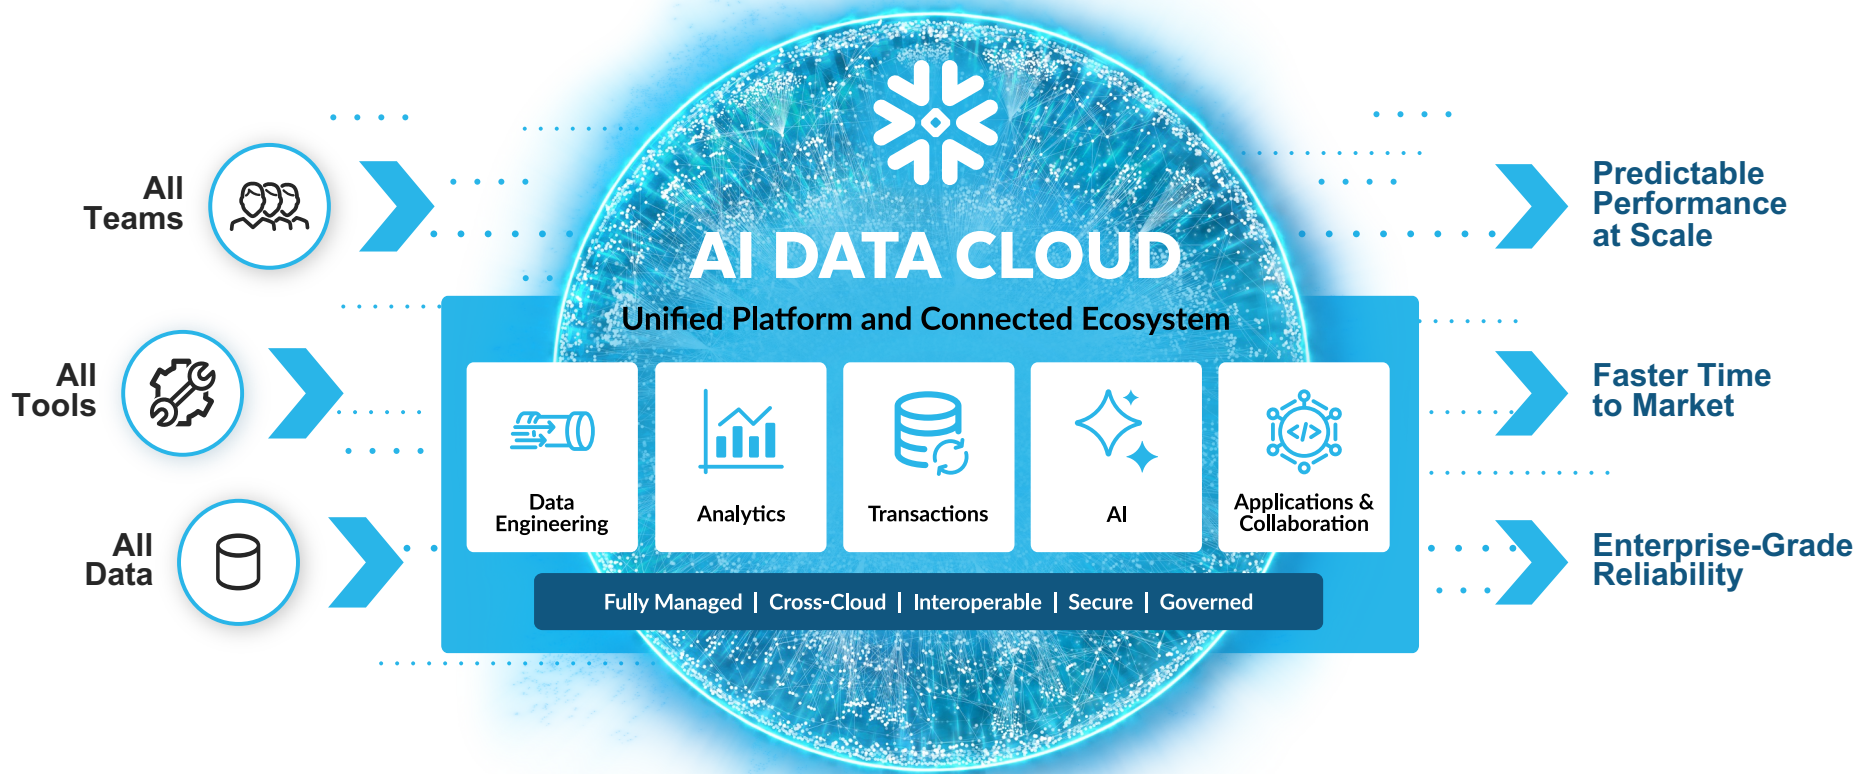
One control plane for the agentic enterprise



THE COMPONENTS OF THE AGENTIC ENTERPRISE



Snowflake powers every data and AI workload



WHAT'S NEW: **AI**

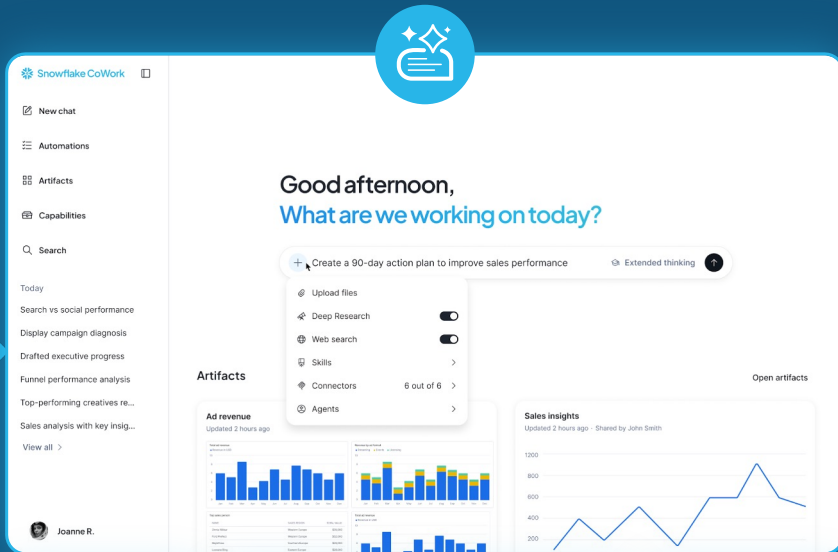


Snowflake CoWork

One AI workspace. Every context. Any outcome. Work gets done.

BRING YOUR CONTEXT

- Email
- Team conversations
- Company data
- Enterprise context
- Apps & MCP

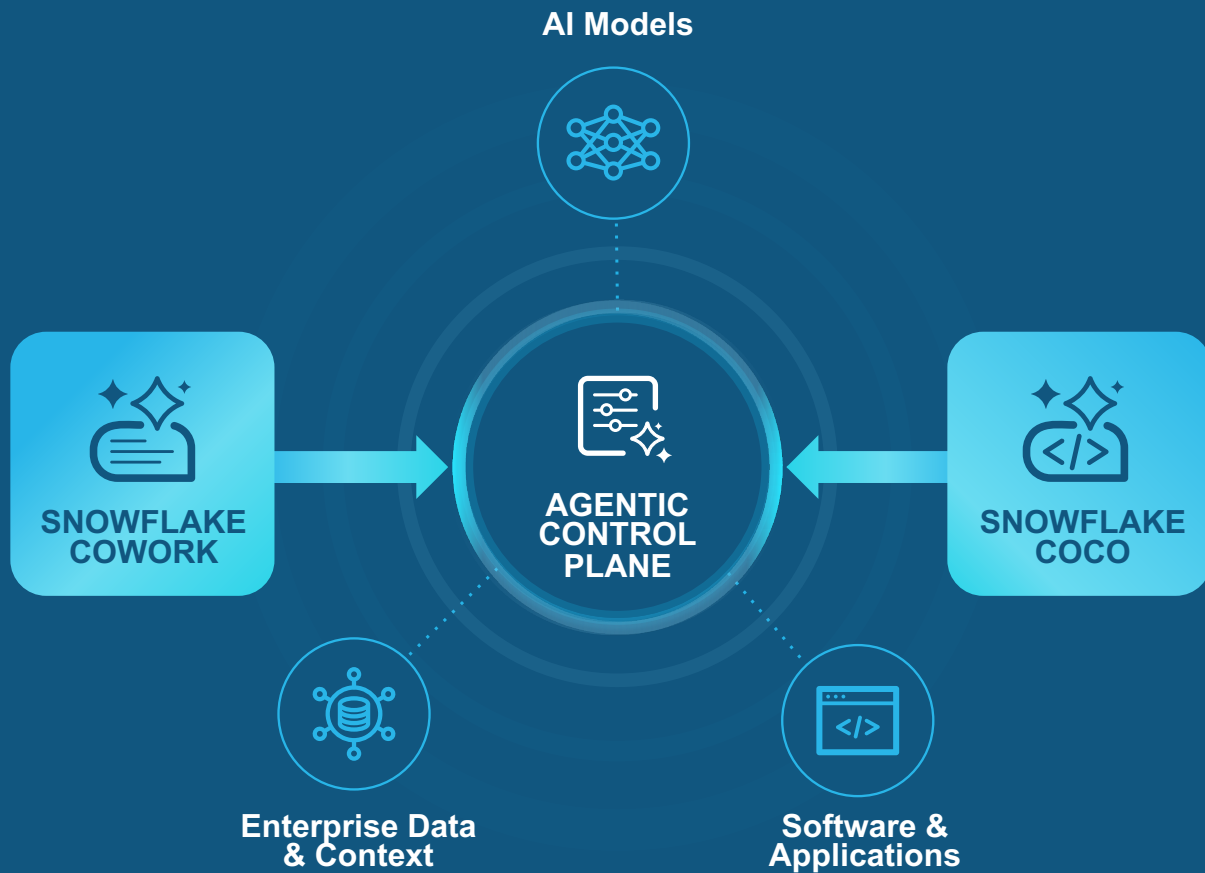


GET WORK DONE

- Presentations
- Documents
- Dashboards
- Emails
- Actions
- Share & Collaborate

Cowork is your Chief of Staff, turning context into clarity, and clarity into action.





Snowflake CoWork

All your knowledge. One trusted Enterprise Agent

From Insight to Action

Personal Work Agent

TRUSTED ENTERPRISE CONTEXT

One governed AI connects teams to the full context of the business

SELF-SERVICE INSIGHTS

Ask questions and get answers without waiting on dashboards or data teams

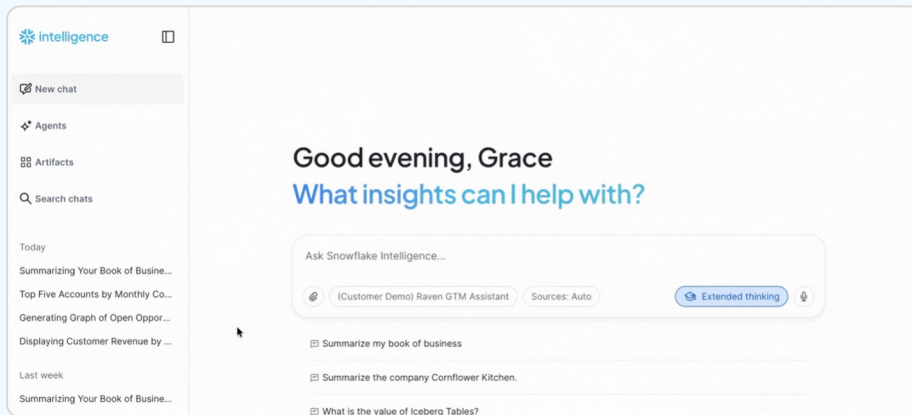
AGENTIC WORK

Reason across enterprise data, automate tasks, and take action in business workflows

The Shift

Conversational Intelligence

Snowflake CoWork agents reason through data to surface what matters, deliver briefs, and recommend actions. It democratizes intelligence.



Snowflake CoWork: What's New

1

Personal Agent

- Knows all enterprise data
- Routes automatically
- Remembers and learns
- Follows you to mobile and existing tools

2

Autonomous Agents

- Proactive checks
- Auto-learns business context
- Produces finished deliverables
- Deep, analytical research

3

Dashboards & Collaboration

- Governed dashboards
- Live, sharable artifacts
- Conversation sharing
- Team knowledge building



Memory, Skills, and Prebuilt Plugins

User Memory

CoWork remembers your preferences, your role, and your recurring tasks, across sessions

PuPr SOON

User Skills

Codify repeatable workflows as a one-line skill, like "pull my pipeline recap", in plain language

PuPr SOON

PLUGINS, Finance & Sales

Prebuilt bundles of skills, business logic, and connectors so your teams aren't starting from zero

PrPr SOON

**Skills are discoverable and shareable across the org.
Every user gets an experience that feels custom-built for their role**



Snowflake CoCo: The AI Coding Agent for Data

Move from 'I have an idea' to 'it's deployed' fast



Accelerates End-to-End Development

Turn data engineering, machine learning, and agent building tasks into simple conversations



Understands Your Enterprise Data Context

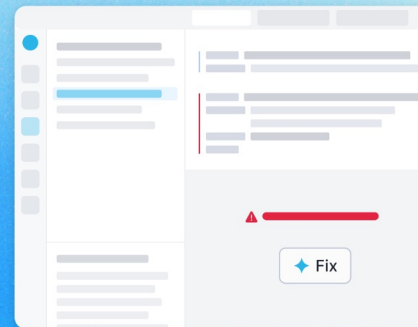
Deeply aware of enterprise data, compute, governance, and operational semantics across the stack



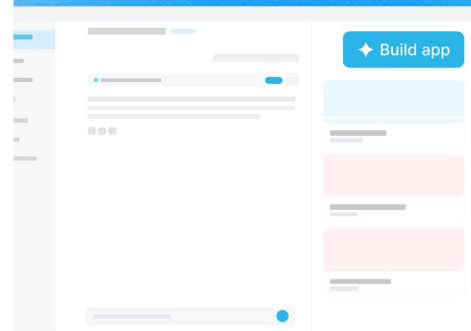
A Governed Agent That Works Wherever You Build

Governed by Snowflake's RBAC and available across Snowsight, your IDE, terminal, and the tools and workflows you already use

Snowsight



Desktop



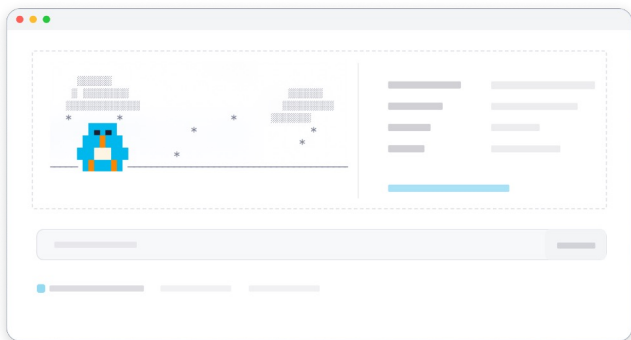
CLI



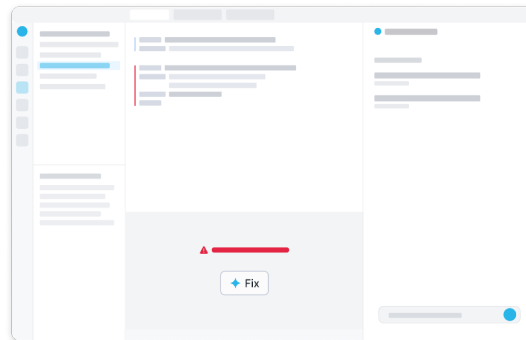
CoCo: The AI Coding Agent for Data



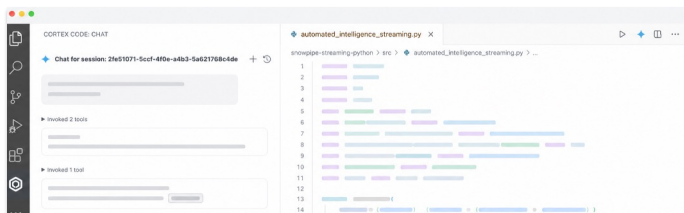
Where you can use CoCo today ...




 **CoCo CLI**



 **CoCo in
Snowsight**



 **CoCo Extension for VS
Code**



 **CoCo Plugin for Claude Code**

Introducing CoCo Desktop

Full agentic development on data, now in a native desktop IDE

One governed surface to build across data stack

Build data pipelines, apps, and agents; debug notebooks, and visualize data flows — all in CoCo

Run agent on schedule, not just on demand

Schedule tasks, and execute multi-step workflows even when you're not at the keyboard.

Extensible by design

MCP integrations and shareable skills/plugins turn best practices into reusable team expertise



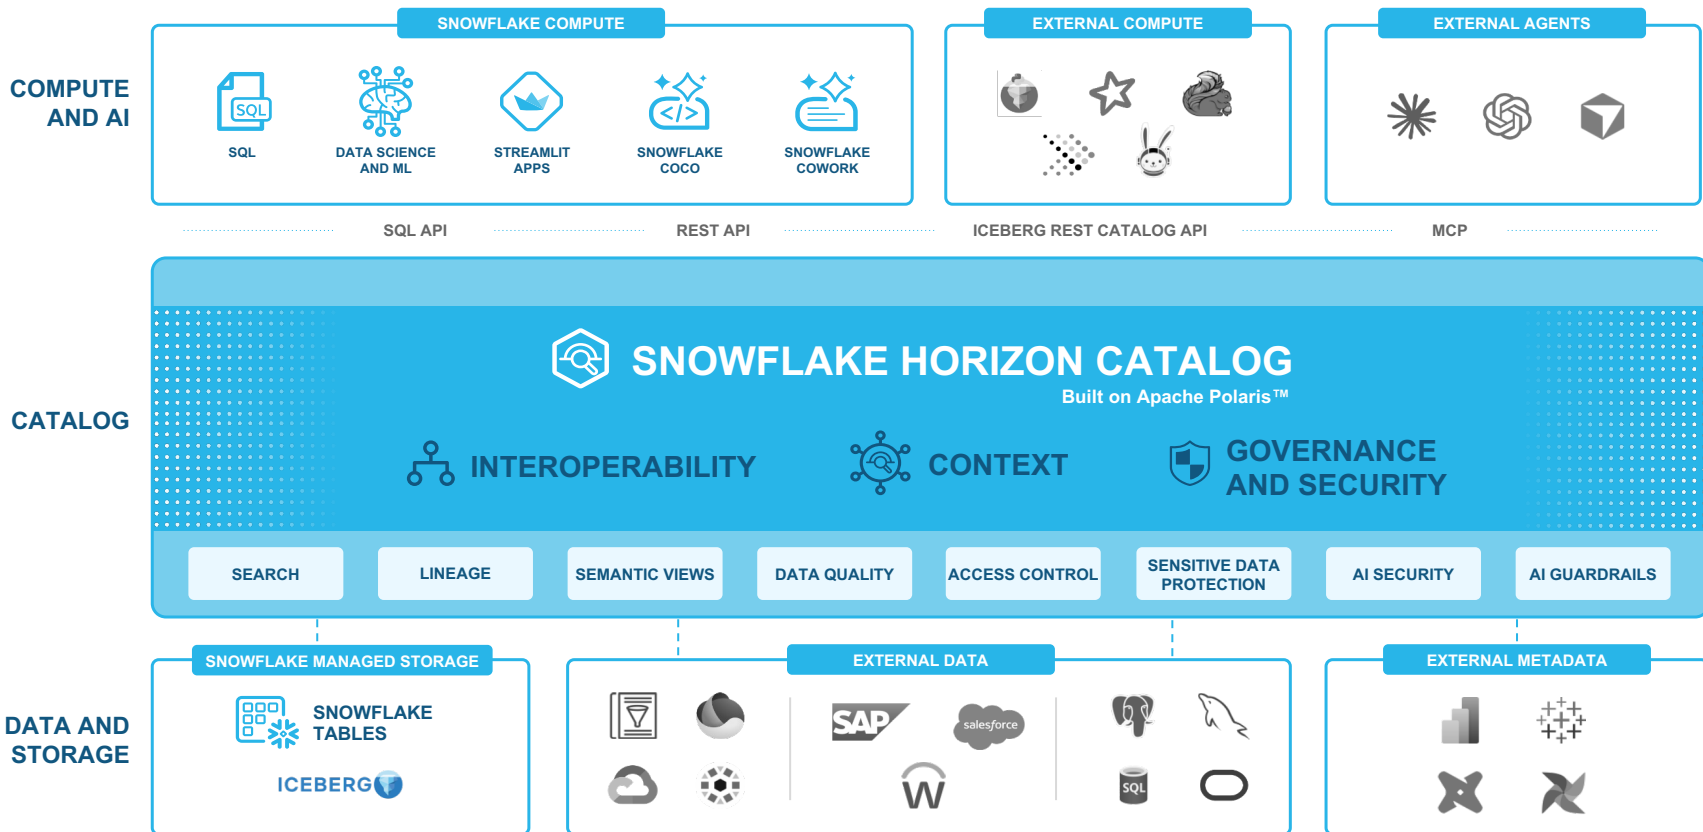
The picture can't be displayed.



WHAT'S NEW: PLATFORM

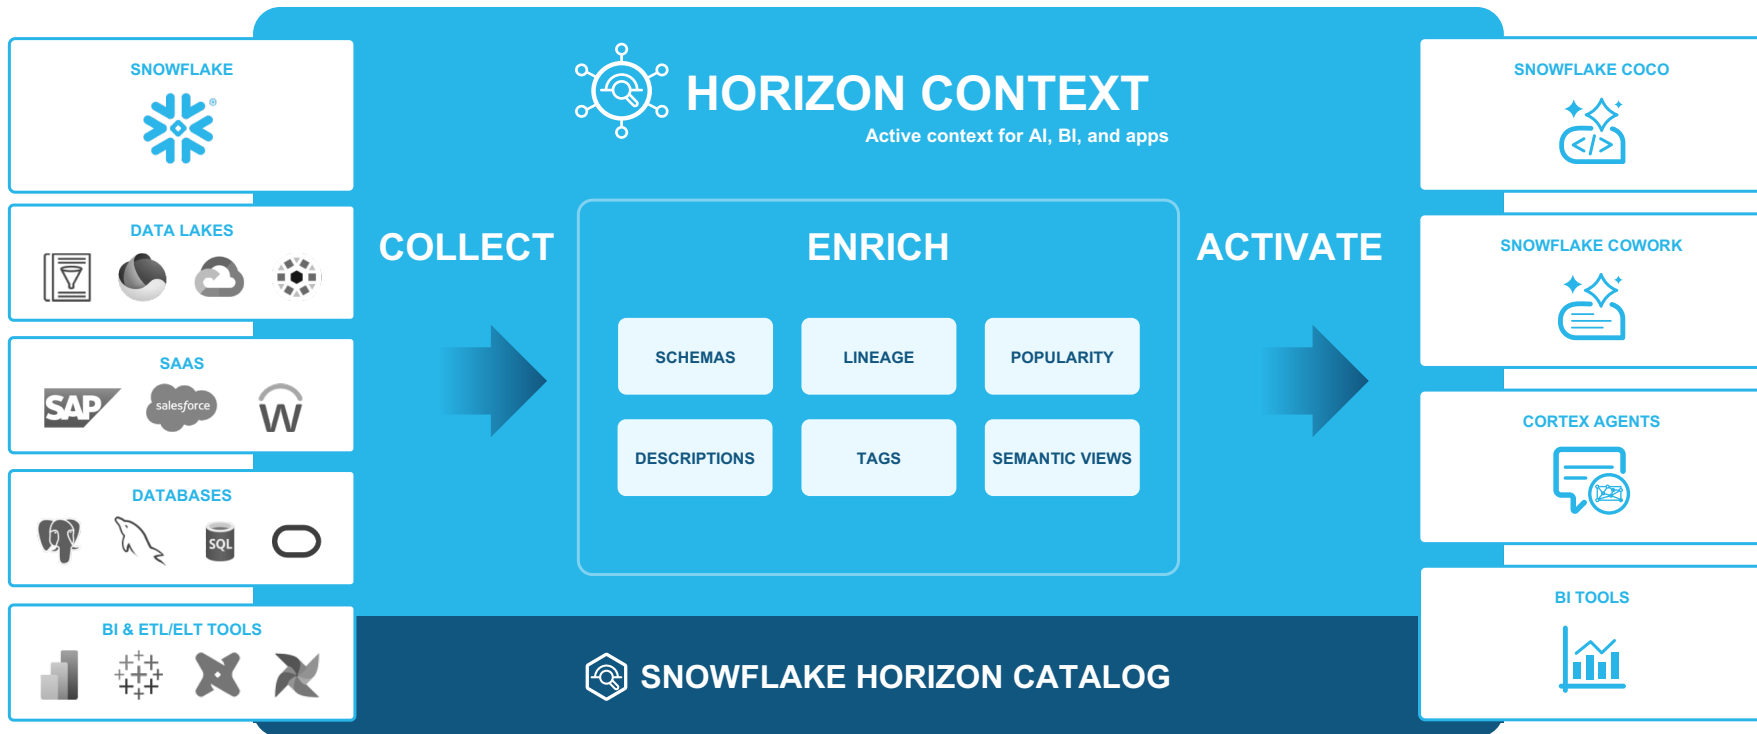


Horizon Catalog: Where Data and AI Connect



Introducing Horizon Context

Active Context for AI, BI, and Apps



Horizon Context Semantic Views

Every AI agent and BI tool reads from the same source

ONE DEFINITION

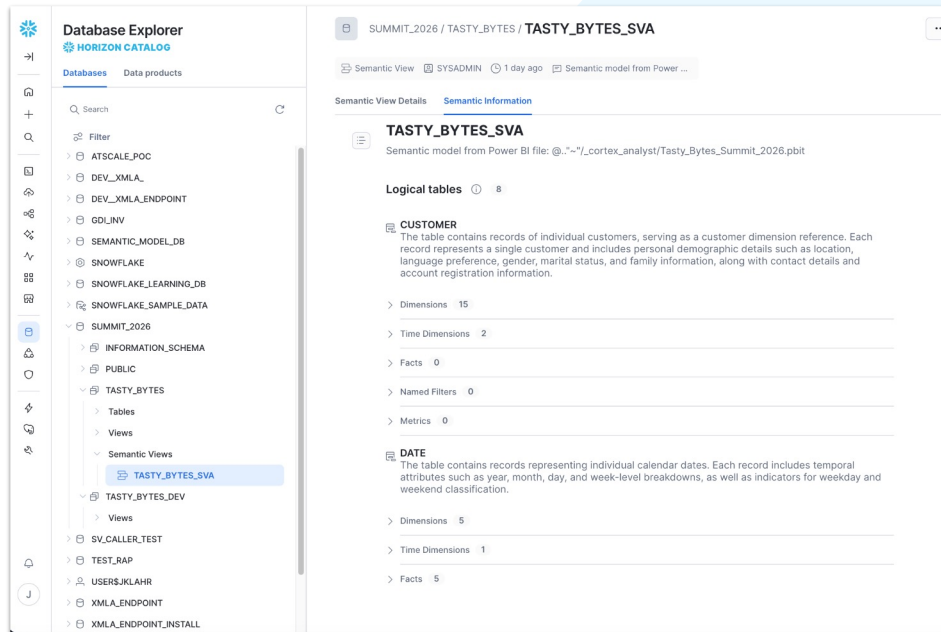
Tables, joins, metrics, filters - described in YAML.
The single source of truth that Cortex Analyst, Power BI, and your custom apps all help resolve against.

NOW MORE POWERFUL

Level of detail calculations, composable definitions you can reuse across models, Type 2 SCD handling, and user-defined materializations with automatic query rewrite.

NOW EASIER

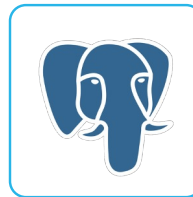
Autopilot generates a complete model from your schema. **Semantic Studio (PrPr)** gives you a full IDE with Git syncing. Cortex Code skills let you iterate in natural language.



Metadata Connectors

Select Star technology, now native to Horizon Catalog

- Connects to external database systems, BI platforms, and ETL/ELT tools
- Crawls external systems and automatically keeps metadata up-to-date
- Mines schemas, dependencies, query/access logs, and more



Additional connectors coming soon



Search Across Your Entire Data Estate

Universal Search, enhanced – plus agentic catalog discovery

WHAT IS IT

Use Universal Search to search for assets in Snowflake and across all your connected sources: PostgreSQL databases, Tableau dashboards, dbt pipelines, and more

WHY USE IT

One place to find anything, built into Snowflake

HOW TO USE IT

Connect your external database, BI, and data pipeline systems, and then just use Universal Search

Type **All**

Tables & views
[View all](#)

FCT_CAMPAIGN_PERFORMANCE

Table | MARKETING > GOLD | updated 5 days ago

Daily campaign metrics including impressions, clicks, spend, conversions, revenue, ...

campaign_sends

PostgreSQL table | marketing > public | updated 13 days ago

Business intelligence
[View all](#)

Campaign Performance Overview

Workbook | Tableau > Marketing | updated 9 days ago

Campaign ROI Scorecard


Dashboard | Power BI > Marketing | updated 2 days ago



WAREHOUSES AND COMPUTE

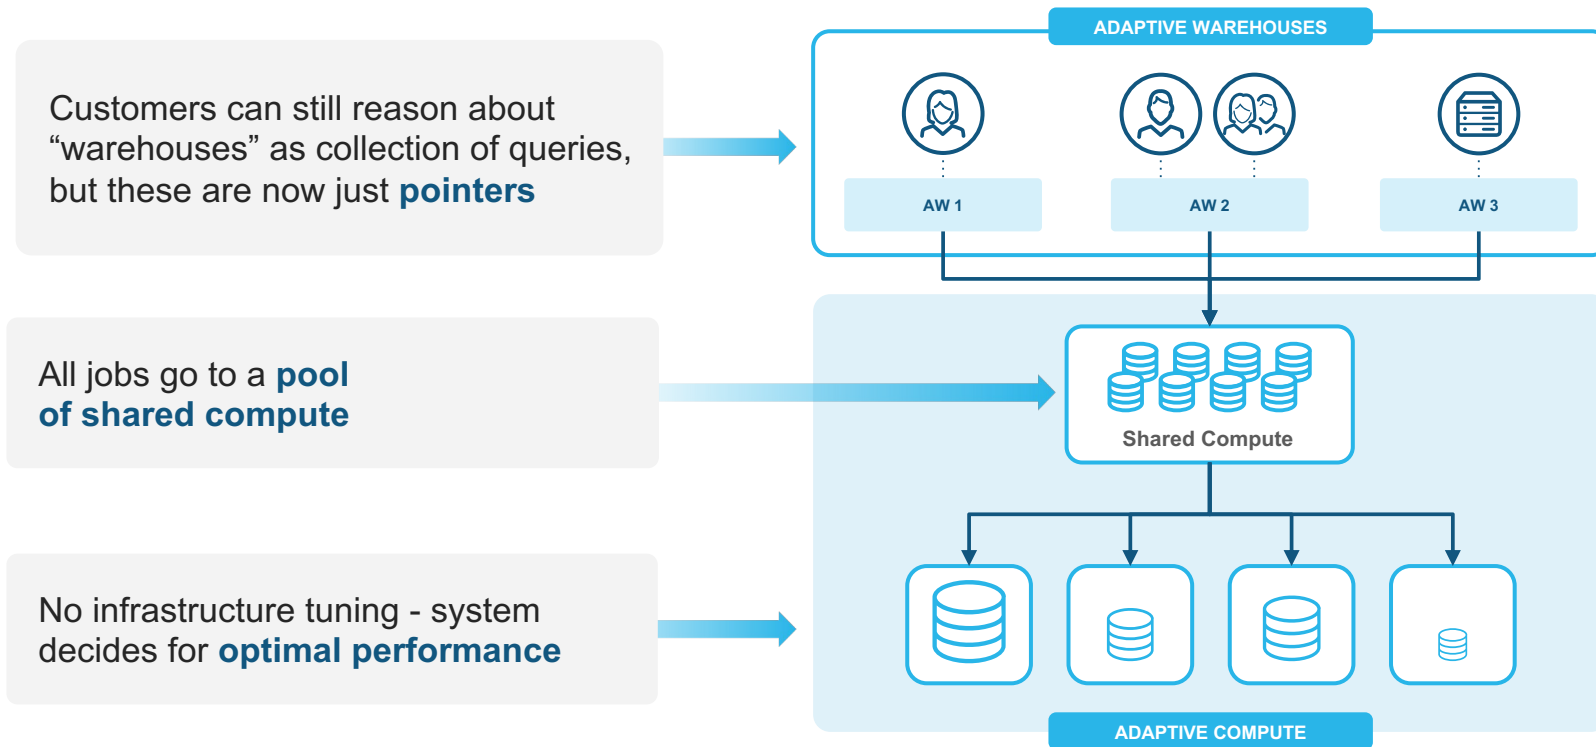


Optimize Every Workload with the Right Compute Resource



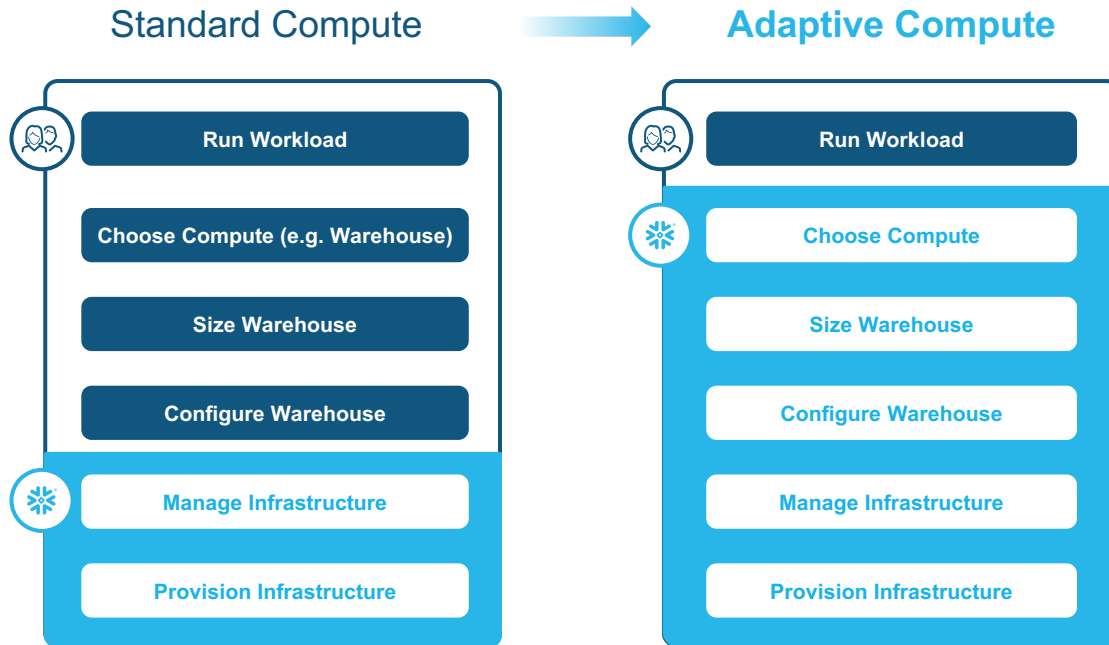
Adaptive Compute

New compute service providing world class performance & effortless operations



Snowflake Compute Evolved

**Scale your
compute
experience**
to make it more
operational
efficient with
better performance



Snowflake Optima

It should be impossible to run inefficiently



Optima Indexing

Automatically identifies recurring point-lookup queries from history and builds hidden search indexes to accelerate them. No DDL, no config, no index billing.



Optima Metadata

Learns “hot” predicate patterns (e.g., common filters in security workloads) and materializes custom pruning metadata to cut I/O for those queries.



Optima Planning

Learns from prior query performance history to optimize future query performance



Optima Clustering

Use workload characteristics to dynamically tune Automatic Clustering for better performance

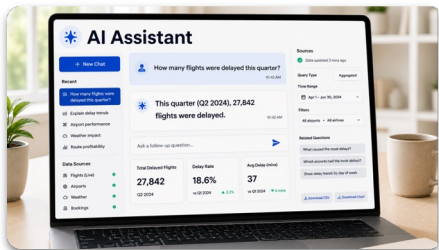
**GA on Gen2 / Adaptive Warehouses at
no additional cost**

**PrPr,
all warehouses**



Interactive Analytics

Built for low-latency, high concurrency workloads



AI & Agents

- Low cost RAGs
- AI Observability
- High concurrency MCPs

Customer-Facing Data Apps

- Data APIs for Apps
- Embedded Analytics
- Dashboards

Operational Analytics

- Trading & Risk Management
- Infra Observability & Alerting
- Supply Chain and Inventory Manager

Low-latency

•

High Concurrency

•

Fresh Data

•

Low Cost per Query

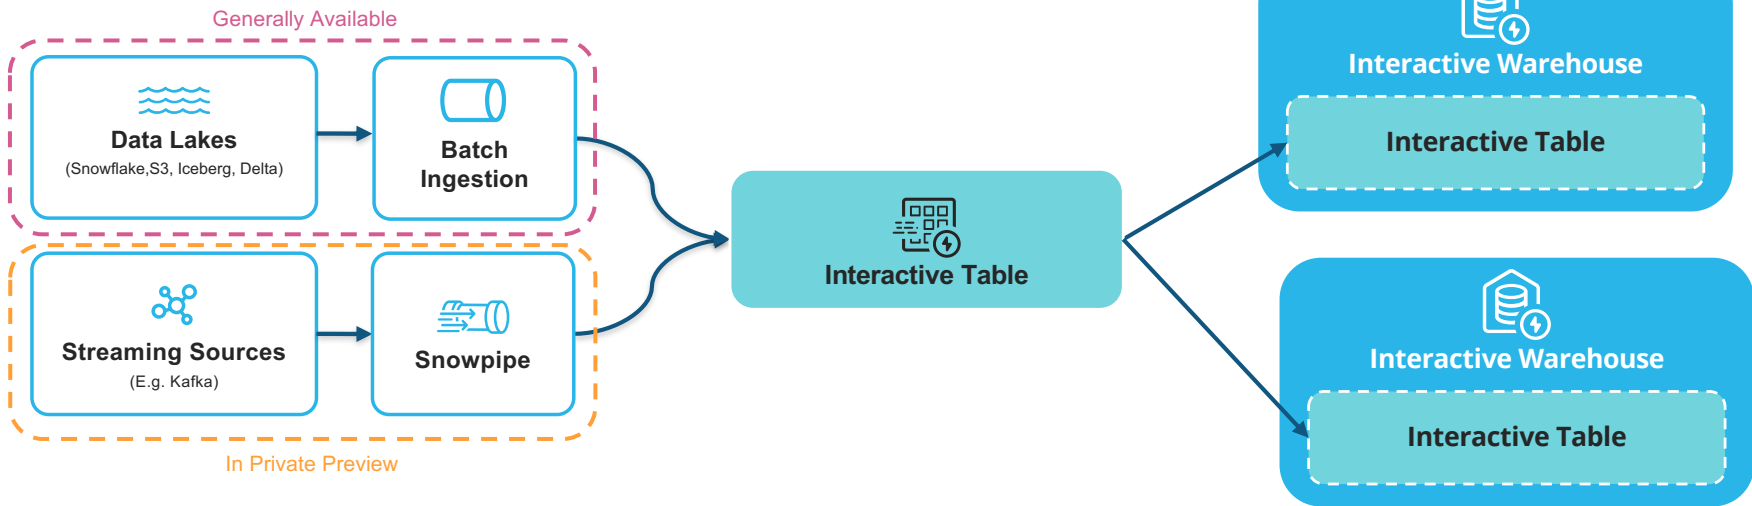


Interactive Compute

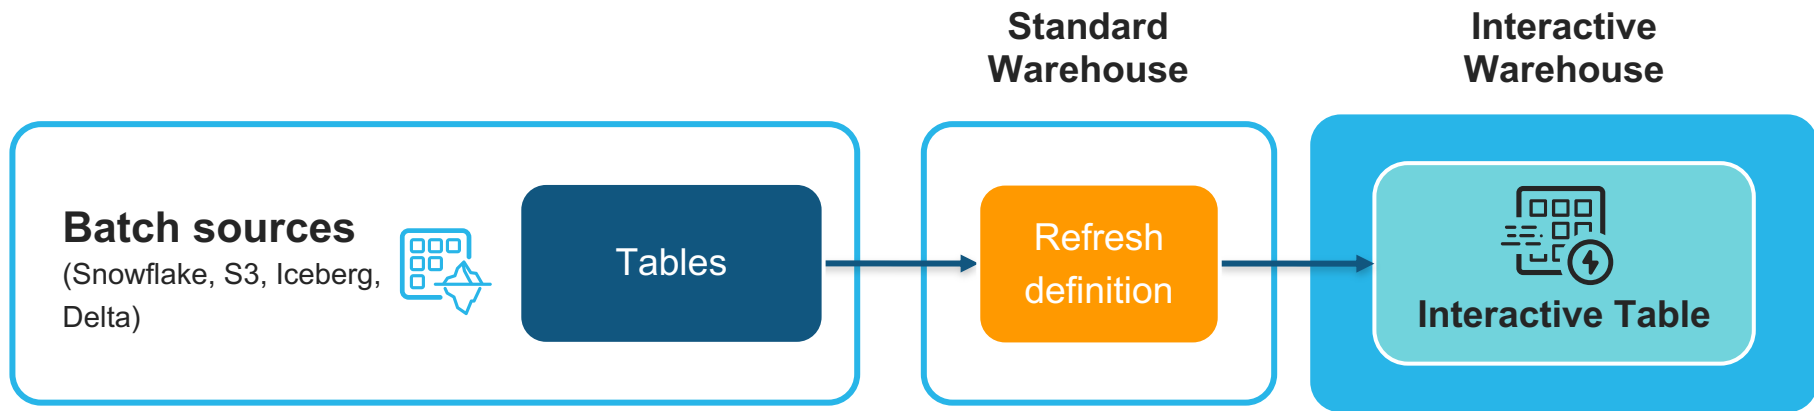
WHAT IS IT

Interactive Table is a new table type that is optimized for real-time streaming ingestion and querying of data.

Interactive Warehouse is a new warehouse that's designed to be always on to serve low-latency queries with pre-warmed caches.



Auto Incremental Refresh (This is how you do DML)

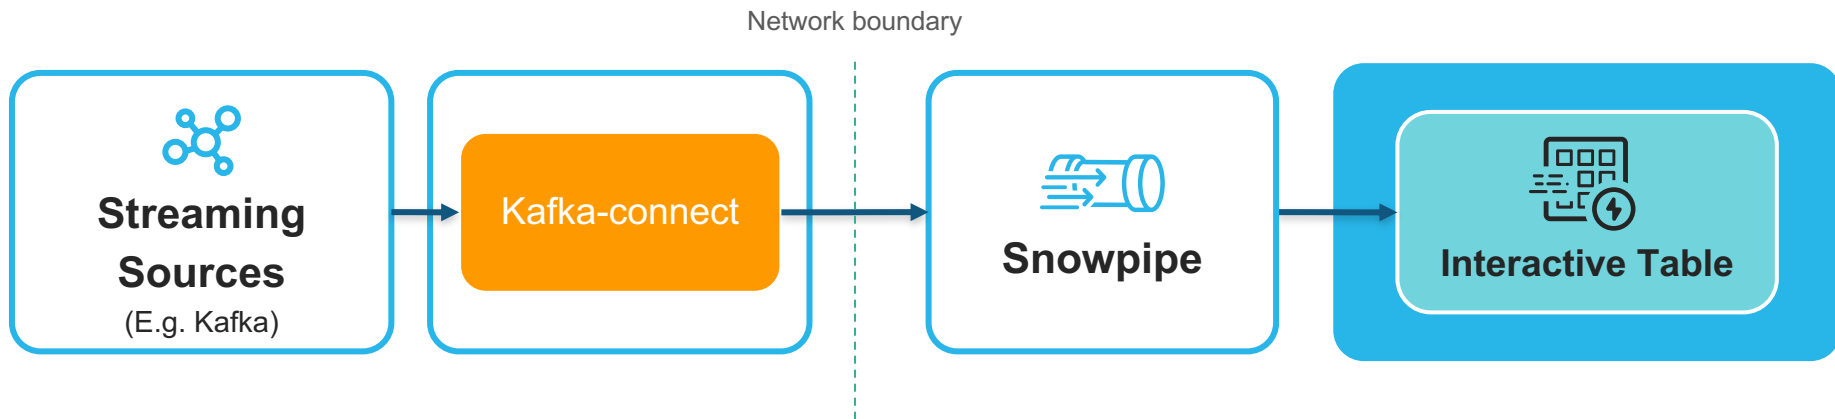


```
CREATE INTERACTIVE TABLE my_dynamic_interactive_table
  CLUSTER BY (c1, c2)
  TARGET_LAG = '5 minutes'*
  WAREHOUSE = my_standard_warehouse ← Refresh warehouse
AS SELECT c1, MAX(c2) FROM my_source_table
GROUP BY c1;
```



Streaming Ingestion (PrPr, all clouds)

Quickstart



- Easy setup, runs regular kafka-connect with Snowpipe v2
- <1s freshness
- 0.01Credit/GB flat rate, no hidden cost (Clustering is included!)
- Yes on schema evolution



Snowflake Compute Evolved

New purpose built compiler and engines

Optimized for your specific workloads
that run better on Snowflake



**Interactive
Compiler**



**Interactive
Engine**



**Unistore
Engine**

New optimizations and engines for analytics in the AI era



Two Paths to Trusted Data, One Powerful Platform



Fully-managed Postgres,
built to deliver the simplest
connection between
OLTP and OLAP



**Single database for analytical
and transactional data,**
with Snowflake security and
governance capabilities baked in



How Snowflake Postgres Connects Your Data

Postgres for Data Mirroring



Postgres to Snowflake

Seamless replication from Postgres to Snowflake in seconds. One command mirrors schemas and tables.

AVAILABLE NOW



Snowflake to Postgres

Seamless replication from Snowflake to Postgres in seconds. One command mirrors schemas and tables.

PREVIEW IN SEPTEMBER

Postgres for Your Data Lake



File Movement

Move files between Postgres and Snowflake, with your internal stage or external object storage as the intermediate layer.

AVAILABLE NOW



Shared Tables

Create Iceberg tables giving your OLTP and OLAP a shared view. Host on internal stage or external object storage.

AVAILABLE NOW



INTEROPERABLE LAKEHOUSE



Interoperability Without Compromise

Architect for agency over your data. Use any engine and accelerate AI on a single data copy.

Capture complete context

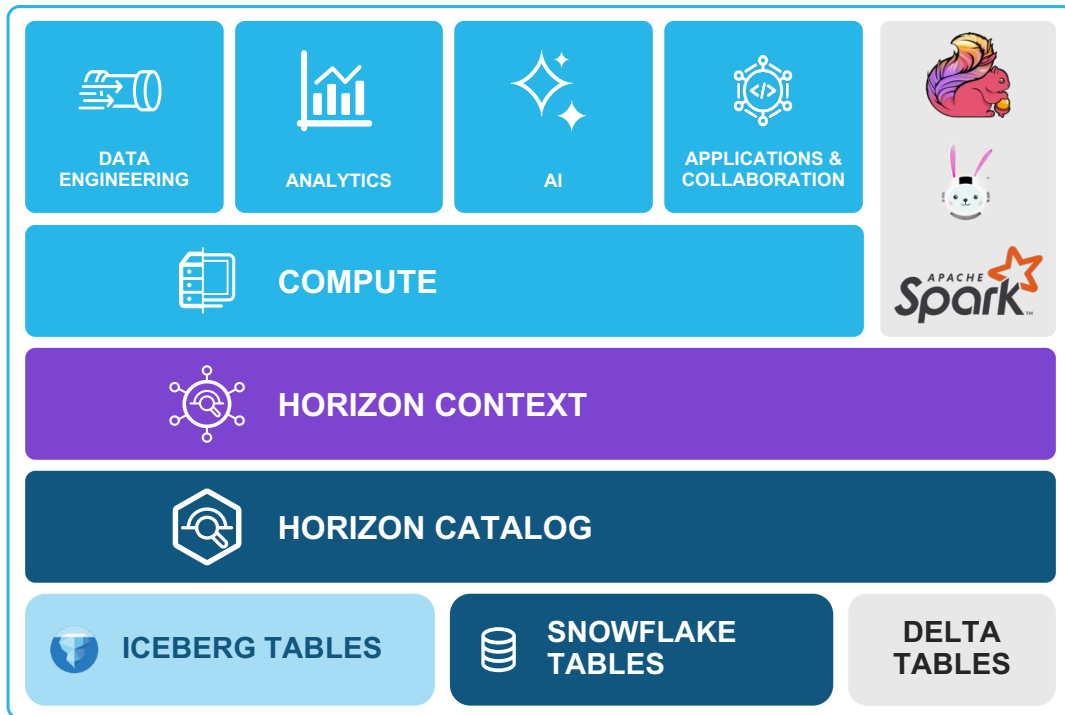
Build toward portable semantic context in your data to increase AI success anywhere

Achieve universal governance

Persist policies across engines to mitigate governance & security risks, without lock-in

Access data in-place

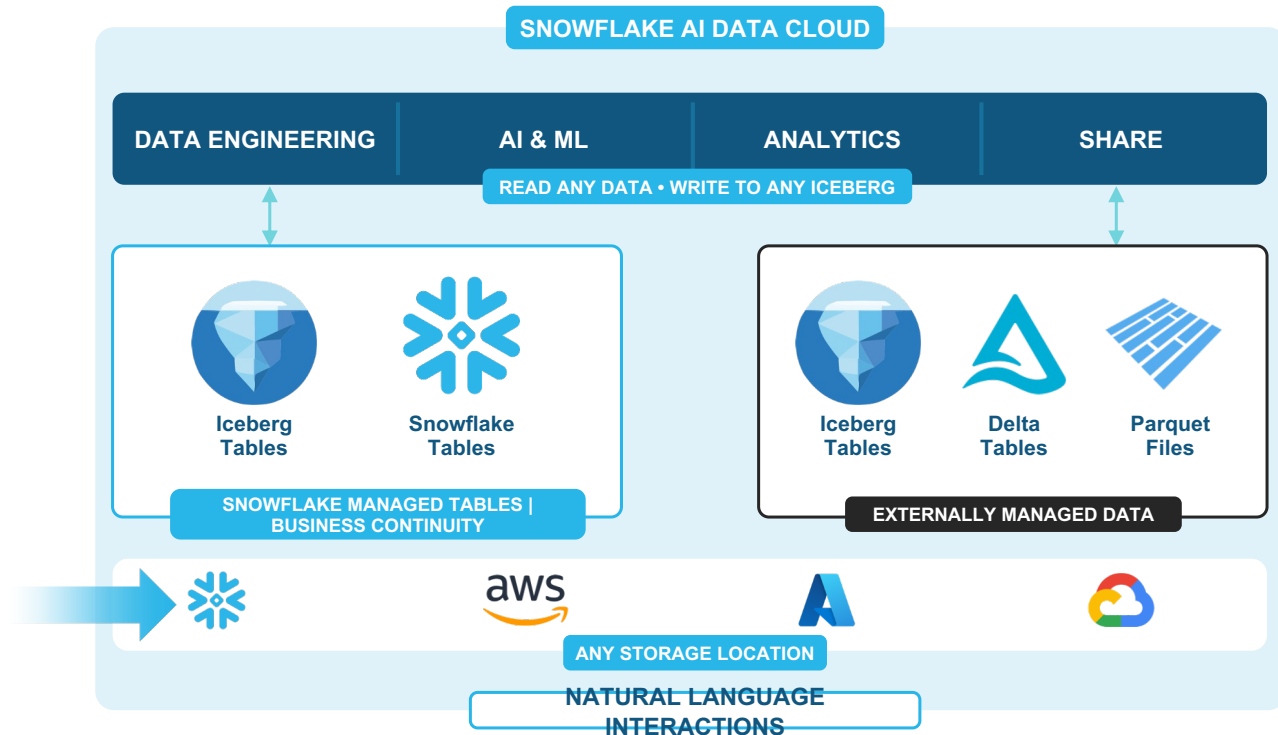
Cut costs with a single governed, resilient & shareable data copy for any operation



Access All Your Data In-Place

Act faster and more efficiently with enterprise performance for **any operation**

- Read and write to any Iceberg table
- Easily deliver business continuity
- Securely share and access live, governed data



Snowflake Storage for Apache Iceberg™

Storage option for open interoperability without the complexity

One Command. Done

Predictable costs and lower effort with a fully managed storage option

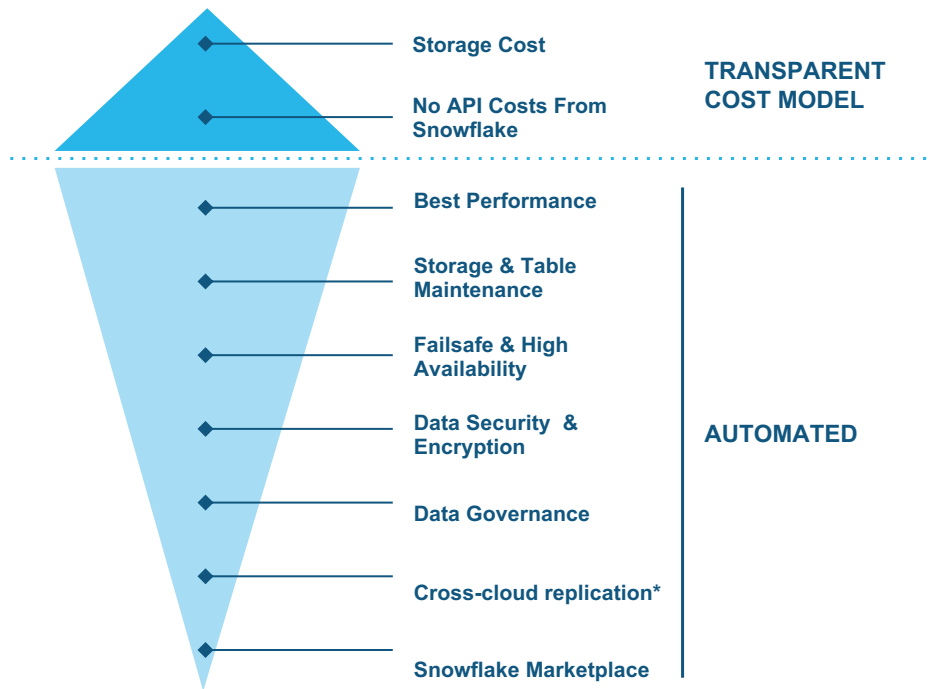
Enterprise-ready by default

Integrated resiliency, private connectivity & native fail-safe recovery

Any engine. One Table.

Peak performance for Snowflake and defaults to boost efficiency on other engines.

```
create iceberg table t (col int)
CATALOG = SNOWFLAKE
EXTERNAL_VOLUME = SNOWFLAKE_MANAGED;
```



Iceberg v3 Improves Semi-structured, Geospatial, and CDC Use Cases

WHAT IS IT

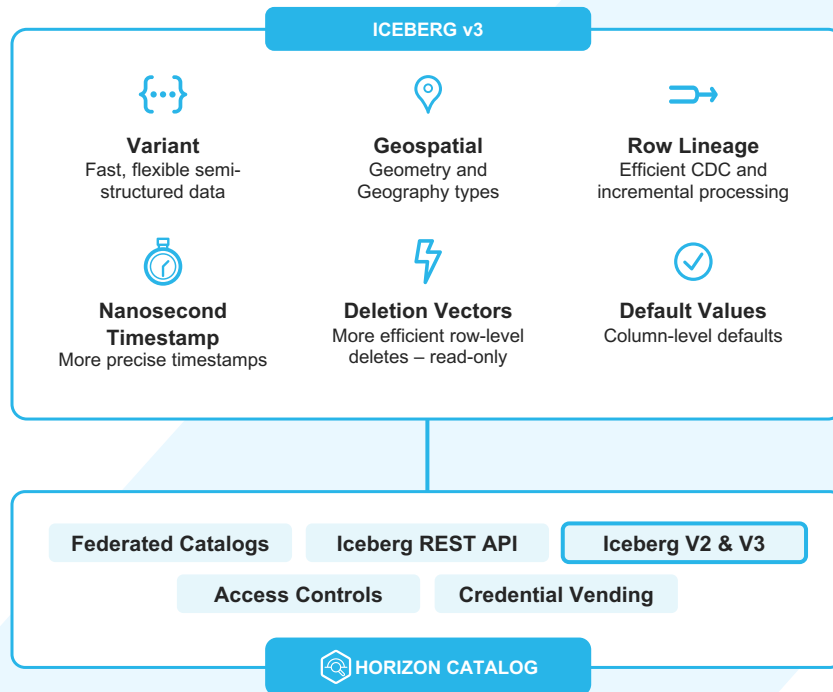
- Snowflake can read and write tables to Horizon Catalog or external catalogs using the Iceberg V3 table spec
- Engines can read V3 tables from Horizon Catalog via Iceberg REST API

WHY USE IT

Better support use cases involving semi-structured data, geospatial data, or CDC that require the interoperability of Iceberg

HOW TO USE IT

Use Horizon Catalog or a catalog integration, define Iceberg tables using the parameter `ICEBERG_VERSION = 3`



Horizon Catalog: One Interoperability Management Plane for All Data

WHAT IS IT

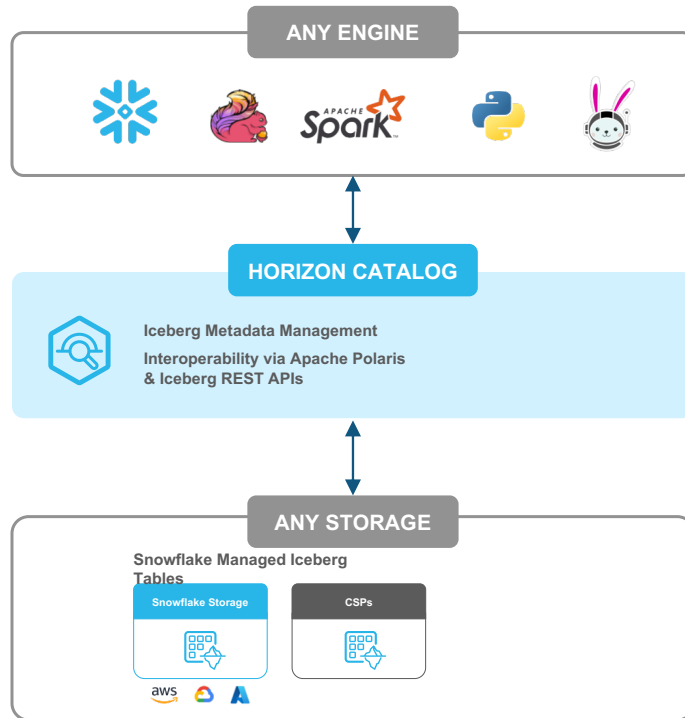
- Manage secure, multi-engine access to Snowflake managed Iceberg tables directly in Horizon Catalog via embedded open APIs [Apache Polaris and Iceberg REST Catalog (IRC)]
- External engines can read from (GA) or write to (GA) Snowflake managed Iceberg tables

WHY USE IT

Interoperability without vendor lock in: Snowflake Horizon Catalog delivers enterprise security, governance, and interoperability without lock-in for your entire data estate

HOW TO USE IT

- Create your Snowflake managed Iceberg tables in Snowflake.
- Manage secure access to Iceberg tables via Horizon Catalog for any compatible external client



External Engine Management for External Iceberg

WHAT IS IT

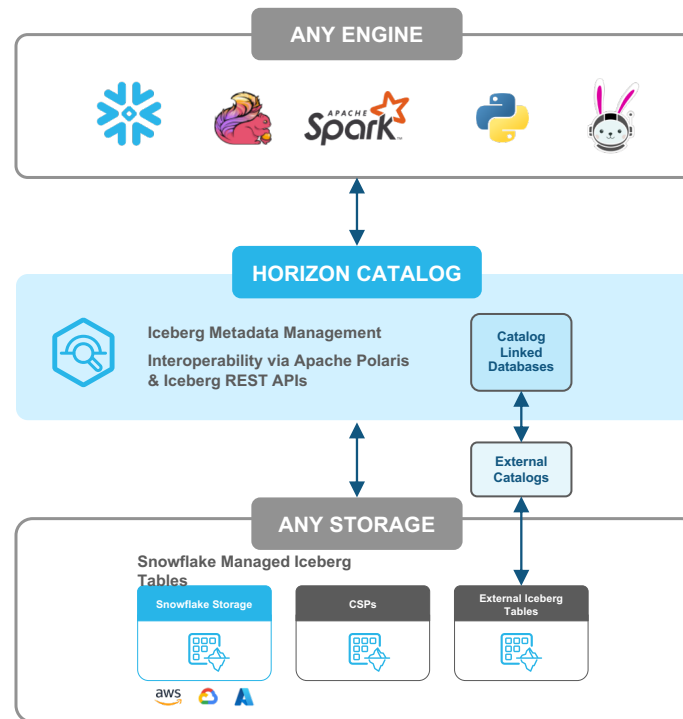
Connected governance for all Iceberg tables. Manage RBAC and security centrally in Horizon, for any Iceberg table, enforcing policies across all external engines.

WHY USE IT

Split governance creates security risks and raises compliance costs. Horizon centralizes Iceberg access across any table on any engine for full read and write operations

HOW TO USE IT

External engines using Horizon Iceberg REST Catalog APIs can now read and write to externally managed tables via Catalog Linked Databases (CLD).



Consistent Policy Enforcement Across Your Ecosystem

WHAT IS IT

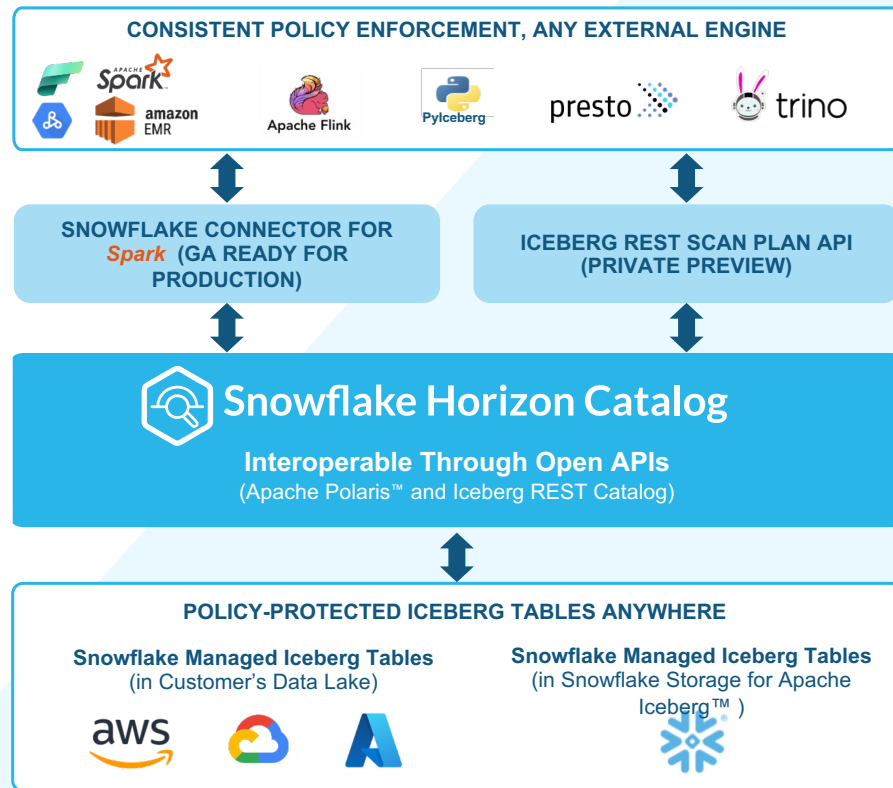
- Enforce row and column masking policies consistently across Iceberg ecosystem through Horizon Catalog
- Works with self-hosted open-source and 3rd-party managed engines
- Allows reads and writes to policy protected Iceberg tables

WHY USE IT

- Promotes consistent row-access and column-masking policies across engines.
- Reduces operational and cost overhead by removing the need for duplicating datasets.
- Balances open-source flexibility with enterprise-grade data protection and trust.

HOW TO USE IT

- Connector (GA): Use the latest Snowflake Connector for Spark package. Run standard Spark queries without changing any existing user code.
- Scan Plan API (Preview): Update your engine to support Iceberg version 1.11.0 and run queries without changing any existing user code.
- Seamless integration with no query changes. Snowflake automatically enforces policies while your external engine handles local compute



Snowflake AIM

AI-driven modernization and virtualization for your data ecosystem on Snowflake



VIRTUALIZATION



RUN TERADATA ON SNOWFLAKE WITHOUT REWRITING

Escape renewal pressure without a risky rewrite

Mitigate rewrite risk — keep applications running

Protect mission-critical operations with zero downtime

Modernize faster and turn migration into AI-driven value



MODERNIZATION



MIGRATE DATA WAREHOUSES TO SNOWFLAKE

Convert code and database objects from leading data warehouses

Support tables, views, stored procedures, and more

End-to-end migration for SQL Server and Amazon Redshift

Automatic translation of legacy ETL workflows into dbt in Snowflake



MIGRATE AND MODERNIZE SPARK WORKLOADS

Analyze python, java, or scala based Spark code to identify APIs and patterns

Assess compatibility with Snowpark and Snowpark Connect

Generate automated conversions to Snowpark equivalents

SNOWFLAKE AIM MIGRATION AGENT

AI-powered code understanding, translation, and migration intelligence across Snowflake AIM modernization workflows

Deploy to Snowflake Native & Managed Iceberg Tables

SUPPORTED SOURCES

ORACLE

SQL SERVER

TERADATA.

AWS

REDSHIFT

GOOGLE
BIGQUERY

GREENPLUM

SYBASE

SYNAPSE

NETEZZA


POSTGRES SQL

DATABRICKS SQL



Snowflake AIM for Virtualization

Exit Teradata fast — with minimal rewrites and zero downtime



KEY BUSINESS OUTCOMES

-  **ESCAPE RENEWAL PRESSURE**
Avoid long-term Teradata contracts and cost escalations
-  **MINIMAL REWRITE RISK**
Run existing applications and tools unchanged
-  **PROTECT OPERATIONS & CONTINUITY**
Zero downtime and business continuity
-  **MODERNIZE ON YOUR TERMS**
Lift and shift now, modernize in phases

Shut down Teradata now. Modernize later. On your timeline.



How Snowflake AIM Orchestrates Migrations

